

Big Data Analytics on Container-Orchestrated Systems

Gerard Casas Saez

University of Colorado Boulder

July 20, 2017

Outline

Introduction & Background

Problem statement

Related work

Approach

Demo

Evaluation & Results

- Scalability

- Maintenance

- Performance

Future work

Introduction & Background

Why?

Data growth

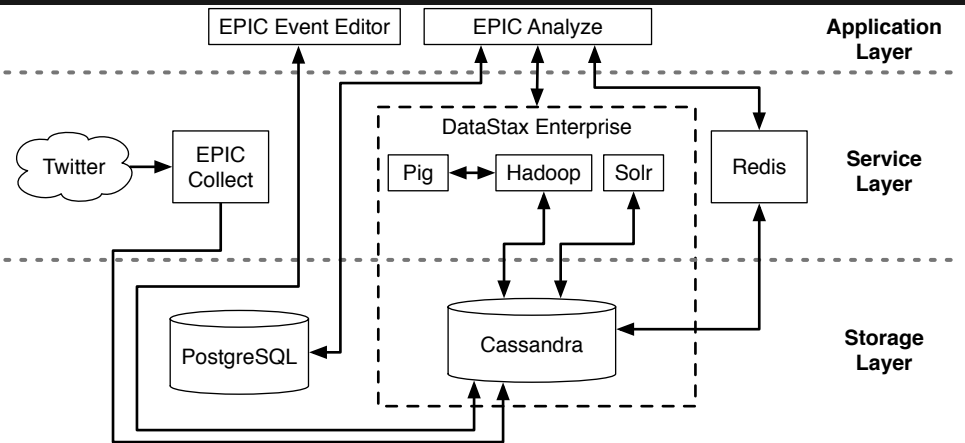
- IOT & Social networks
- Increase internet traffic: 3x internet traffic by 2021
- Scale up Big Data Analytics System
- Keeping maintenance at low cost
- Container-orchestrated make infrastructure easier



Keep up with data growth

Background: Project EPIC

- EPIC Collect
- EPIC Analyze



Background: Containerization

- Operating-system-level virtualization
- Use host machine system resources
- **Docker** most used alternative
- Development microservices

Background:

Container-orchestration systems

- Container interaction abstraction
- Great to deploy microservices architectures
- Apache Mesos vs **Kubernetes**



Background: Microservices Architecture

- Small & specific
- Better scalability
- Loosely-coupled & highly-cohesive
- Orchestration <> **Coreography**

Background: Coreography microservice architecture

- Easier to extend
- Asynchronous
- PubSub interaction
- Messaging system: **Apache Kafka**



Problem statement

Problem statement

1. Advantages and/or limitations from existing infrastructure
 - 1.1 More reliable?
 - 1.2 More scalable?
2. Lower maintenance costs than the existing infrastructure?
 - 2.1 Easier to deploy?
 - 2.2 Easier to upgrade?
 - 2.3 More resilient to failures?

Related work

- SMACK: Spark, Mesos, Akka, Cassandra and Kafka [Raul Estrada et al. 2016]
- Hadoop ecosystem [Han Hu et al. 2014]
- Lambda architecture [Zirije Hasani et al. 2014]

Approach

Features

- Event management
- Real-time collection of streaming Twitter data
- Real-time classification of incoming tweets
- Data Analysis


Custom components

- Event Manager
- Infrastructure Controller
- Twitter tracker
- Twitter Normalizer

Cassandra table structure

```
CREATE TABLE twitter_analytics.tweet (  
  id uuid, t_id text, event_kw text, event_name text, hashtags list<text>,   
  media_url text, t_coordinates text, t_created_at timestamp,   
  t_favorite_count int, t_favorited boolean, t_geo text,   
  t_is_a_retweet boolean, t_lang text, t_retweet_count int,   
  t_retweeted boolean, t_text text, u_created_at timestamp,   
  u_description text, u_favourites_count int, u_followers_count int,   
  u_friends_count int, u_geo_enabled boolean, u_id text, u_lang text,   
  u_listed_count int, u_location text, u_name text, u_screen_name text,   
  u_statuses_count int, u_time_zone text, u_url text, u_utc_offset int,   
  um_id text, um_name text, um_screen_name text, urls list<text>,   
  PRIMARY KEY (id, t_id))
```

Listing 1: Tweets CQL table script



Demo time!

Let's track an event...

Event Manager UI

...and analyze it!

Zeppelin Notebook

Smiley

DB code	smiley
Tokens	:D, :, XD, LOL, WTF, XP, :', :-

Hearth

DB code	hearth
Tokens	<3, love

nba-draft

DB code	nba-draft
Tokens	nba, draft, Lonzo, Lavar, basketball, nba draft

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
cassandra	cassandra-0	2/2	Running	1	2d
cassandra	cassandra-1	2/2	Running	2	11d
cassandra	cassandra-2	2/2	Running	7	22d
cassandra	spark-master-controller-bkd17	1/1	Running	0	22d
default	hearth-event-parser-2160998245-m19st	1/1	Running	9	2d
default	k8s-controller-3919038388-75tkk	1/1	Running	0	2d
default	smiley-event-parser-3033807940-c0cwj	1/1	Running	9	2d
default	twitter-tracker-2482383360-s0kz1	1/1	Running	5	2d
frontend	eventmanager-ui-3464180876-h605f	1/1	Running	0	23d
frontend	zeppelin-3633522582-t0kng	1/1	Running	0	2d
kafka	kafka-0	1/1	Running	3	11d
kafka	kafka-1	1/1	Running	0	2d
kafka	zoo-0	1/1	Running	0	2d
kafka	zoo-1	1/1	Running	0	2d
kafka	zoo-2	1/1	Running	0	25d
kube-system	heapster-v1.3.0-4211727876-kv0cl	2/2	Running	0	11d
kube-system	kube-dns-806549836-43lvx	3/3	Running	0	11d
kube-system	kube-dns-autoscaler-2528518105-2wlsr	1/1	Running	1	27d
kube-system	kube-proxy-gke-development-development-dcfa2eb3-2jhh	1/1	Running	0	2d
kube-system	kube-proxy-gke-development-development-dcfa2eb3-l5xb	1/1	Running	1	26d

cassandraspark-samples/analytics

Real time twitter Analytics

Press run above to update stats

Columns: id, t_id, event_kw, event_name, hashtags, media_url, t_coordinates, t_created_at, t_favorite_count, t_favorited, t_geo, t_is_a_retweet, t_lang, t_retweet_count, t_retweeted, t_text, u_created_at, u_description, u_favourites_count, u_followers_count, u_friends_count, u_geo_enabled, u_id, u_lang, u_listed_count, u_location, u_name, u_screen_name, u_statuses_count, u_time_zone, u_url, u_utc_offset, um_id, um_name, um_screen_name, urls

SQL Context configured


Stats on tweets

Different stats on all the current dataset available on the database


Total tweet count

```
import org.apache.spark._
import com.datastax.spark.connector._
import com.datastax.spark.connector._
table: com.datastax.spark.connector.rdd.CassandraTableScanRDD[com.datastax.spark.connector.CassandraRow] = CassandraTableScanRDD[0] at RDD at CassandraRDD.scala:15
res0: IndexedSeq[String] = WrappedArray(id, t_id, event_kw, event_name, hashtags, media_url, t_coordinates, t_created_at, t_favorite_count, t_favorited, t_geo, t_is_a_retweet, t_lang, t_retweet_count, t_retweeted, t_text, u_created_at, u_description, u_favourites_count, u_followers_count, u_friends_count, u_geo_enabled, u_id, u_lang, u_listed_count, u_location, u_name, u_screen_name, u_statuses_count, u_time_zone, u_url, u_utc_offset, um_id, um_name, um_screen_name, urls)
35713760
```


Tweets by event



Most influential accounts (by followers)



Languages used in all tweets



Evaluation & Results

Reliability

Current vs Prototype

- Threads to monitoring
- Every minute script check
- Check log file size
- Analyze manual
- Kubernetes abstraction
- Auto recovery
- Rolling update
- Node assignation depending on resources usage

Scalability

Current vs Prototype

- Monolith
- Manual process
- Shared session state
- Load balancing
- Kubernetes replica specification
- Stateless microservices
- Independently scalable
- Auto scale

Maintenance

Current infrastructure

- Large & complex
- Strong use of frameworks
- Manual deployment
- Dedicated machines



Prototype

- Easier to maintain: less code, faster development
- Technology flexibility
- Easier to deploy: YAML description files
- Flexibility on cloud provider
- Available built-in tools



Code lines

EPIC Analyze 5086 lines

Twitter tracker 108 lines

Kubernetes controller 145 lines

Event manager 2090 lines

Tweet normalizer 209 lines

Performance

WordCount


```
val table = sc.cassandraTable("twitter_analytics", "tweet")
val words = table.select("t_text").flatMap(l => l.getString("t_text").split(" "))
    .map(word => (word.toLowerCase,1)).reduceByKey(_ + _).map(_._2.swap).
    sortByKey(false,1)
```

Listing 2: WordCount Spark script

```
(1) ShuffledRDD[112] at sortByKey at <console>:31 []
+- (176) MapPartitionsRDD[111] at map at <console>:31 []
    | ShuffledRDD[110] at reduceByKey at <console>:31 []
+- (176) MapPartitionsRDD[109] at map at <console>:31 []
    | MapPartitionsRDD[108] at flatMap at <console>:31 []
    | CassandraTableScanRDD[107] at RDD at CassandraRDD.scala:15 []
```

Listing 3: Debug string rdd

WordCount



Number of tweets	Time (s)	Tweets/sec
490199	238	2060
1400884	469	2987
21680851	1107	19585
27199614	1500	18133
54395957	4228	12866

Future work

- Improve resource specification (CPU, memory)
- Better Cassandra structure
- Unified authentication
- Extend the system with new features

Questions?

Thank you!

Twitter @casassaez

Website gerard.space

Repo github.com/casassg/thesis